

参考文献

REFERENCE

- [1] YANG H, LIN Z, WANG W, et al. Memory³: Language modeling with explicit memory[J]. Journal of Machine Learning, 2024, 3(3): 300-346.
- [2] LI Z, SONG S, XI C, et al. Memos: A memory os for ai system[J]. arXiv preprint arXiv:2507.03724, 2025.
- [3] VASWANI A, SHAZEER N, PARMAR N, et al. Attention is all you need[C]//Advances in Neural Information Processing Systems: volume 30, 2017: 5998-6008.
- [4] LIU J, SHEN Z, HE Y, et al. Towards out-of-distribution generalization: A survey[J]. arXiv preprint arXiv:2108.13624, 2021.
- [5] GURURANGAN S, MARASOVIĆ A, SWAYAMDIPTA S, et al. Don't stop pretraining: Adapt language models to domains and tasks[J]. arXiv preprint arXiv:2004.10964, 2020.
- [6] OUYANG L, WU J, JIANG X, et al. Training language models to follow instructions with human feedback[C]//Advances in Neural Information Processing Systems (NeurIPS). 2022.
- [7] RAFAILOV R, SHARMA A, MITCHELL E, et al. Direct preference optimization: Your language model is secretly a reward model[J]. Advances in neural information processing systems, 2023, 36: 53728-53741.
- [8] GAO Y, XIONG Y, GAO X, et al. Retrieval-augmented generation for large language models: A survey[J]. arXiv preprint arXiv:2312.10997, 2023, 2(1).
- [9] YAN S Q, GU J C, ZHU Y, et al. Corrective retrieval augmented generation[J]. 2024.
- [10] MADAAN A, TANDON N, GUPTA P, et al. Self-refine: Iterative refinement with self-feedback[J]. Advances in Neural Information Processing Systems, 2023, 36: 46534-46594.
- [11] OpenAI. Prompt engineering[EB/OL]. . <https://platform.openai.com/docs/guides/prompt-engineering>.
- [12] OpenAI. Best practices for prompt engineering with the OpenAI API[EB/OL]. <https://help.openai.com/en/articles/6654000-best-practices-for-prompt-engineering-with-the-openai-api>.
- [13] Anthropic. Effective context engineering for AI agents[EB/OL]. 2025. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>.
- [14] WENG L. Llm powered autonomous agents[EB/OL]. 2023. <https://lilianweng.github.io/posts/2023-06-23-agent/>.
- [15] DU Y, HUANG W, ZHENG D, et al. Rethinking memory in llm based agents: Representations, operations, and emerging topics[Z]. 2025.
- [16] SHAN L, LUO S, ZHU Z, et al. Cognitive memory in large language models[J]. arXiv preprint arXiv:2504.02441, 2025.
- [17] MALKOV Y A, YASHUNIN D A. Efficient and robust approximate nearest neighbor search

- using hierarchical navigable small world graphs[J]. *IEEE transactions on pattern analysis and machine intelligence*, 2018, 42(4): 824-836.
- [18] ATKINSON R C, SHIFFRIN R M. Human memory: A proposed system and its control processes[M]//*Psychology of learning and motivation: volume 2*. Elsevier, 1968: 89-195.
 - [19] DOYLE J. A truth maintenance system[J]. *Artificial intelligence*, 1979, 12(3): 231-272.
 - [20] PACKER C, FANG V, PATIL S, et al. Memgpt: Towards llms as operating systems.[J]. 2023.
 - [21] ASAI A, WU Z, WANG Y, et al. Self-rag: Learning to retrieve, generate, and critique through self-reflection[J]. 2024.
 - [22] GUU K, LEE K, TUNG Z, et al. Retrieval augmented language model pre-training[C]//*International conference on machine learning*. PMLR, 2020: 3929-3938.
 - [23] LEWIS P, PEREZ E, PIKTUS A, et al. Retrieval-augmented generation for knowledge-intensive nlp tasks[J]. *Advances in neural information processing systems*, 2020, 33: 9459-9474.
 - [24] KIRKPATRICK J, PASCANU R, RABINOWITZ N, et al. Overcoming catastrophic forgetting in neural networks[J]. *Proceedings of the national academy of sciences*, 2017, 114(13): 3521-3526.
 - [25] PARK J S, O'BRIEN J, CAI C J, et al. Generative agents: Interactive simulacra of human behavior[C/OL]//*UIST '23: Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*. New York, NY, USA: Association for Computing Machinery, 2023. DOI: 10.1145/3586183.3606763.
 - [26] WANG W, DONG L, CHENG H, et al. Augmenting language models with long-term memory[J]. *Advances in Neural Information Processing Systems*, 2023, 36: 74530-74543.
 - [27] ZHONG W, GUO L, GAO Q, et al. Memorybank: Enhancing large language models with long-term memory[C]//*Proceedings of the AAAI Conference on Artificial Intelligence: volume 38*. 2024: 19724-19731.
 - [28] WANG R Y, STRONG D M. Beyond accuracy: What data quality means to data consumers[J]. *Journal of management information systems*, 1996, 12(4): 5-33.
 - [29] WAND Y, WANG R Y. Anchoring data quality dimensions in ontological foundations[J]. *Communications of the ACM*, 1996, 39(11): 86-95.
 - [30] MAHARANA A, LEE D H, TULYAKOV S, et al. Evaluating very long-term conversational memory of LLM agents[C/OL]//*KU L W, MARTINS A, SRIKUMAR V. Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Bangkok, Thailand: Association for Computational Linguistics, 2024: 13851-13870. DOI: 10.18653/v1/2024.acl-long.747.
 - [31] CHHIKARA P, KHANT D, ARYAN S, et al. Mem0: Building production-ready ai agents with scalable long-term memory[J]. *arXiv preprint arXiv:2504.19413*, 2025.
 - [32] HINDER F, VAQUET V, HAMMER B. One or two things we know about concept drift—a survey on monitoring in evolving environments. part a: detecting concept drift[J]. *Frontiers in Artificial Intelligence*, 2024, 7: 1330257.
 - [33] WILIE B, CAHYAWIJAYA S, ISHII E, et al. Belief revision: The adaptability of large language models reasoning[J]. *arXiv preprint arXiv:2406.19764*, 2024.
 - [34] SMUHA N A. Regulation 2024/1689 of the eur. parl. & council of june 13, 2024 (eu artificial intelligence act)[J]. *International Legal Materials*, 2025: 1-148.
 - [35] LEE M, KIM K, KIM T, et al. Selective generation for controllable language models[J]. *Advances in Neural Information Processing Systems*, 2024, 37: 50494-50527.
 - [36] HELLAND P. Data on the outside versus data on the inside[J]. *Communications of the ACM*, 2020, 63(11): 111-118.
 - [37] ZHANG Z, DAI Q, BO X, et al. A survey on the memory mechanism of large language model-based agents[J]. *ACM Transactions on Information Systems*, 2025, 43(6): 1-47.
 - [38] YU H, CHEN T, FENG J, et al. Memagent: Reshaping long-context llm with multi-convrl-based memory agent[Z]. 2025.
 - [39] YAO Y, WANG P, TIAN B, et al. Editing Large Language Models: Problems, Methods, and

- Opportunities[M/OL]. arXiv, 2023. DOI: 10.48550/arXiv.2305.13172.
- [40] MENG K, BAU D, ANDONIAN A J, et al. Locating and editing factual associations in GPT[C/OL]//Advances in Neural Information Processing Systems. 2022. DOI: 10.48550/arXiv.2202.05262.
 - [41] MENG K, SHARMA A S, ANDONIAN A, et al. Mass-Editing Memory in a Transformer[M/OL]. arXiv, 2023. DOI: 10.48550/arXiv.2210.07229.
 - [42] CAO N D, AZIZ W, TITOV I. Editing Factual Knowledge in Language Models[M/OL]. arXiv, 2021. DOI: 10.48550/arXiv.2104.08164.
 - [43] MITCHELL E, LIN C, BOSSELUT A, et al. Fast Model Editing at Scale[C]//2021.
 - [44] TAN C, ZHANG G, FU J. Massive Editing for Large Language Models via Meta Learning[C]//2023.
 - [45] MITCHELL E, LIN C, BOSSELUT A, et al. Memory-Based Model Editing at Scale[C]//Proceedings of the 39th International Conference on Machine Learning. PMLR, 2022: 15817-15831.
 - [46] HARTVIGSEN T, SANKARANARAYANAN S, PALANGI H, et al. Aging with GRACE: Lifelong Model Editing with Discrete Key-Value Adaptors[M/OL]. arXiv, 2023. DOI: 10.48550/arXiv.2211.11031.
 - [47] LIU N F, LIN K, HEWITT J, et al. Lost in the middle: How language models use long contexts[J]. Transactions of the Association for Computational Linguistics (TACL), 2024.
 - [48] OUYANG S, YAN J, HSU I, et al. Reasoningbank: Scaling agent self-evolving with reasoning memory[J]. arXiv preprint arXiv:2509.25140, 2025.
 - [49] YAN S, YANG X, HUANG Z, et al. Memory-r1: Enhancing large language model agents to manage and utilize memories via reinforcement learning[Z]. 2025.
 - [50] SONG M, ZHENG M. A survey of query optimization in large language models[J]. arXiv preprint arXiv:2412.17558, 2024.
 - [51] GAO L, MA X, LIN J, et al. Precise zero-shot dense retrieval without relevance labels[C]//Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). 2023: 1762-1777.
 - [52] LOUIS M, FORMAL T, DEJEAN H, et al. OSCAR: Online soft compression and reranking[J]. arXiv preprint arXiv:2504.07109, 2025.
 - [53] LIU L, YANG X, SHEN Y, et al. Think-in-memory: Recalling and post-thinking enable llms with long-term memory[Z]. 2023.
 - [54] XU W, LIANG Z, MEI K, et al. A-mem: Agentic memory for llm agents[J]. arXiv preprint arXiv:2502.12110, 2025.
 - [55] SHINN N, CASSANO F, GOPINATH A, et al. Reflexion: Language agents with verbal reinforcement learning[J]. Advances in Neural Information Processing Systems, 2023, 36: 8634-8652.
 - [56] ZHAO A, HUANG D, XU Q, et al. Expel: Llm agents are experiential learners[C]//Proceedings of the AAAI Conference on Artificial Intelligence: volume 38. 2024: 19632-19642.
 - [57] FANG R, LIANG Y, WANG X, et al. Memp: Exploring agent procedural memory[Z]. 2025.
 - [58] KANG J, JI M, ZHAO Z, et al. Memory os of ai agent[Z]. 2025.
 - [59] WANG Y, TAKANOBU R, LIANG Z, et al. Mem- α : Learning memory construction via reinforcement learning[Z]. 2025.
 - [60] PETRONI F, ROCKTÄSCHEL T, MILLER A H, et al. Language models as knowledge bases?[C]//In: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2019. 2019.
 - [61] GUPTA A, RAO A, ANUMANCHIPALLI G. Model editing at scale leads to gradual and catastrophic forgetting[J]. arXiv preprint arXiv:2401.07453, 2024.
 - [62] COHEN R, BIRAN E, YORAN O, et al. Evaluating the ripple effects of knowledge editing in language models[J/OL]. Transactions of the Association for Computational Linguistics, 2024, 12: 283-298. DOI: 10.1162/tacl_a_00644.

- [63] LI Y, et al. Does knowledge localization hold true? Surprising differences between entity and relation perspectives in language models[J]. 2024.
- [64] ZHANG N, YAO Y, TIAN B, et al. A comprehensive study of knowledge editing for large language models[J]. arXiv preprint arXiv:2401.01286, 2024.
- [65] HUANG L, YU W, MA W, et al. A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions[Z]. 2023.
- [66] WALLAT J, HEUSS M, DE RIJKE M, et al. Correctness is not faithfulness in RAG attributions[Z]. 2024.
- [67] XU R, QI Z, GUO Z, et al. Knowledge conflicts for LLMs: A survey[C]//Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing. 2024: 8541-8565.
- [68] SHI W, HAN X, LEWIS M, et al. Trusting your evidence: Hallucinate less with context-aware decoding[C]//Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers). 2024: 783-791.
- [69] ZHANG Q, XIANG Z, XIAO Y, et al. FaithfulRAG: Fact-level conflict modeling for context-faithful retrieval-augmented generation[C]//Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). 2025.
- [70] TAN Z, YAN J, HSU I H, et al. In prospect and retrospect: Reflective memory management for long-term personalized dialogue agents[C]//Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). 2025: 8416-8439.
- [71] HONG K, TROYNIKOV A, HUBER J. Context rot: How increasing input tokens impacts llm performance[R]. Chroma, 2025.
- [72] YAO S, ZHAO J, YU D, et al. React: Synergizing reasoning and acting in language models[C]//The Eleventh International Conference on Learning Representations. 2023.
- [73] JIMENEZ GUTIERREZ B, SHU Y, GU Y, et al. Hipporag: Neurobiologically inspired long-term memory for large language models[J]. Advances in Neural Information Processing Systems, 2024, 37: 59532-59569.
- [74] LI Z, JIN X, LI W, et al. Temporal knowledge graph reasoning based on evolutionary representation learning[C/OL]//SIGIR '21: Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval. 2021: 408-417. DOI: 10.1145/3404835.3462963.
- [75] JIN W, QU M, JIN X, et al. Recurrent event network: Autoregressive structure inference over temporal knowledge graphs[C]//Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP). 2020: 6669-6683.
- [76] TULVING E. Episodic and semantic memory[M]//TULVING E, DONALDSON W. Organization of Memory. New York: Academic Press, 1972: 381-403.
- [77] LOGAN J. Continuum memory architectures for long-horizon llm agents[Z]. 2025.
- [78] SQUIRE L R, ALVAREZ P. Retrograde amnesia and memory consolidation: A neurobiological perspective[J/OL]. Current Opinion in Neurobiology, 1995, 5(2): 169-177. DOI: 10.1016/0959-4388(95)80023-9.
- [79] JIANG X, LI F, ZHAO H, et al. Long term memory: The foundation of ai self-evolution[Z]. 2024.
- [80] SARTHI P, ABDULLAH S, TULI A, et al. Raptor: Recursive abstractive processing for tree-organized retrieval[C]//The Twelfth International Conference on Learning Representations. 2024.
- [81] CHEN W, QI G, LI W, et al. Decide then retrieve: A training-free framework with uncertainty-guided triggering and dual-path retrieval[Z]. 2026.
- [82] CHEN D, NIU S, LI K, et al. Halumem: Evaluating hallucinations in memory systems of agents[Z]. 2026.
- [83] HU Y, WANG Y, MCAULEY J. Evaluating memory in LLM agents via incremental multi-turn interactions[C]//The Fourteenth International Conference on Learning Representations.

- 2026.
- [84] EBBINGHAUS H. Über das gedächtnis: Untersuchungen zur experimentellen psychologie[M]. Leipzig: Duncker & Humblot, 1885.
 - [85] XIONG Z, LIN Y, XIE W, et al. How memory management impacts llm agents: An empirical study of experience-following behavior[Z]. 2025.
 - [86] LIU J, SU Y, XIA P, et al. Simplemem: Efficient lifelong memory for llm agents[J]. arXiv preprint arXiv:2601.02553, 2025.
 - [87] YU Y, YAO L, XIE Y, et al. Agentic memory: Learning unified long-term and short-term memory management for large language model agents[J]. arXiv preprint arXiv:2601.01885, 2026.
 - [88] WU Y, LIANG S, ZHANG C, et al. From human memory to ai memory: A survey on memory mechanisms in the era of llms[J]. arXiv preprint arXiv:2504.15965, 2025.
 - [89] LangChain Inc. Langgraph: Building stateful multi-actor applications with llms[EB/OL]. 2024. <https://www.langchain.com/langgraph>.
 - [90] QIAN C, XIE Z, WANG Y, et al. Scaling large language model-based multi-agent collaboration[J]. arXiv preprint arXiv:2406.07155, 2024.
 - [91] WANG G, XIE Y, JIANG Y, et al. Voyager: An open-ended embodied agent with large language models[J]. arXiv preprint arXiv:2305.16291, 2023.
 - [92] WANG Y, CHEN X. Mirix: Multi-agent memory system for llm-based agents[J]. arXiv preprint arXiv:2507.07957, 2025.
 - [93] ZHANG G, FU M, WANG K, et al. G-memory: Tracing hierarchical memory for multi-agent systems[C]//The Thirty-ninth Annual Conference on Neural Information Processing Systems. 2025.
 - [94] QIAN C, DANG Y, LI J, et al. Experiential co-learning of software-developing agents[C]//Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). 2024: 5628-5640.
 - [95] WU S, SHU K. Memory in llm-based multi-agent systems: Mechanisms, challenges, and collective intelligence[J/OL]. TechRxiv preprint, 2025. DOI: 10.36227/techrxiv.176539617.79044553/v1.
 - [96] TRAN K T, DAO D, NGUYEN M D, et al. Multi-agent collaboration with large language models: A survey[J]. arXiv preprint arXiv:2501.06322, 2025.
 - [97] NII H P. The blackboard model of problem solving and the evolution of blackboard architectures[J/OL]. AI Magazine, 1986, 7(2): 38-53. DOI: 10.1609/aimag.v7i2.537.
 - [98] SAGIROVA A, KURATOV Y, BURTSEV M. Srmt: Shared memory for multi-agent lifelong pathfinding[J]. arXiv preprint arXiv:2501.13200, 2025.
 - [99] LIU M, ROTTA T, SOUZA V, et al. Designing multi-agent intelligence[EB/OL]. 2025. <https://developer.microsoft.com/blog/designing-multi-agent-intelligence>.
 - [100] WU Q, BANSAL G, ZHANG J, et al. Autogen: Enabling next-gen LLM applications via multi-agent conversations[C]//First Conference on Language Modeling. 2024.
 - [101] SAKS J. Memory for multi agent systems[EB/OL]. 2024. <https://medium.com/@juliansaks/memory-for-multi-agent-systems-b0f63ec87a51>.
 - [102] WANG L, MA C, FENG X, et al. A survey on large language model based autonomous agents[J/OL]. Frontiers of Computer Science, 2024, 18(6). DOI: 10.1007/s11704-024-40231-1.
 - [103] XI Z, CHEN W, GUO X, et al. The rise and potential of large language model based agents: a survey[J/OL]. Science China Information Sciences, 2025, 68(2): 121101. DOI: 10.1007/s11432-024-4222-0.
 - [104] XU B, PENG Z, LEI B, et al. Rewoo: Decoupling reasoning from observations for efficient augmented language models[Z]. 2023.
 - [105] WANG L, XU W, LAN Y, et al. Plan-and-solve prompting: Improving zero-shot chain-of-thought reasoning by large language models[C/OL]//ROGERS A, BOYD-GRABER J, OKAZAKI N. Proceedings of the 61st Annual Meeting of the Association for Computa-

- tional Linguistics (Volume 1: Long Papers). Toronto, Canada: Association for Computational Linguistics, 2023: 2609-2634. DOI: 10.18653/v1/2023.acl-long.147.
- [106] YAO S, YU D, ZHAO J, et al. Tree of thoughts: Deliberate problem solving with large language models[C]//OH A, NAUMANN T, GLOBERSON A, et al. Advances in Neural Information Processing Systems: volume 36. Curran Associates, Inc., 2023: 11809-11822.
 - [107] BESTA M, BLACH N, KUBICEK A, et al. Graph of thoughts: solving elaborate problems with large language models[C/OL]//AAAI'24/IAAI'24/EAAI'24: Proceedings of the 38th AAAI Conference on Artificial Intelligence and 36th Conference on Innovative Applications of Artificial Intelligence and 14th Symposium on Educational Advances in Artificial Intelligence. AAAI Press, 2024. DOI: 10.1609/aaai.v38i16.29720.
 - [108] HUANG X, LIU W, CHEN X, et al. Understanding the planning of llm agents: A survey[Z]. 2024.
 - [109] KIM Y, GU K, PARK C, et al. Towards a science of scaling agent systems[Z]. 2025.
 - [110] GOU Z, SHAO Z, GONG Y, et al. CRITIC: Large language models can self-correct with tool-interactive critiquing[C]//The Twelfth International Conference on Learning Representations. 2024.
 - [111] YIN X, WANG X, PAN L, et al. Gödel agent: A self-referential agent framework for recursively self-improvement[C/OL]//CHE W, NABENDE J, SHUTOVA E, et al. Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). Vienna, Austria: Association for Computational Linguistics, 2025: 27890-27913. DOI: 10.18653/v1/2025.acl-long.1354.
 - [112] ZHOU P, PUJARA J, REN X, et al. Self-discover: Large language models self-compose reasoning structures[C/OL]//GLOBERSON A, MACKEY L, BELGRAVE D, et al. Advances in Neural Information Processing Systems: volume 37. Curran Associates, Inc., 2024: 126032-126058. DOI: 10.52202/079017-4004.
 - [113] KIM S, MOON S, TABRIZI R, et al. An LLM compiler for parallel function calling[C]//Forty-first International Conference on Machine Learning. 2024.
 - [114] HE G, DEMARTINI G, GADIRAJU U. Plan-then-execute: An empirical study of user trust and team performance when using llm agents as a daily assistant[C]//2025 CHI Conference on Human Factors in Computing Systems, CHI 2025. ACM, 2025: 414.
 - [115] XIE Y, KAWAGUCHI K, ZHAO Y, et al. Self-evaluation guided beam search for reasoning[C]//OH A, NAUMANN T, GLOBERSON A, et al. Advances in Neural Information Processing Systems: volume 36. Curran Associates, Inc., 2023: 41618-41650.
 - [116] HAO S, GU Y, MA H, et al. Reasoning with language model is planning with world model[C/OL]//BOUAMOR H, PINO J, BALI K. Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing. Singapore: Association for Computational Linguistics, 2023: 8154-8173. DOI: 10.18653/v1/2023.emnlp-main.507.
 - [117] ZHOU A, YAN K, SHLAPENTOKH-ROTHMAN M, et al. Language agent tree search unifies reasoning, acting, and planning in language models[C]//ICML'24: Proceedings of the 41st International Conference on Machine Learning. Vienna, Austria: JMLR.org, 2024.
 - [118] QIAO S, ZHANG N, FANG R, et al. AutoAct: Automatic agent learning from scratch for QA via self-planning[C/OL]//KU L W, MARTINS A, SRIKUMAR V. Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). Bangkok, Thailand: Association for Computational Linguistics, 2024: 3003-3021. DOI: 10.18653/v1/2024.acl-long.165.
 - [119] CHENG Y, ZHANG C, ZHANG Z, et al. Exploring large language model based intelligent agents: Definitions, methods, and prospects[Z]. 2024.
 - [120] LUO J, ZHANG W, YUAN Y, et al. Large language model agent: A survey on methodology, applications and challenges[Z]. 2025.
 - [121] BOIKO D A, MACKNIGHT R, KLINE B, et al. Autonomous chemical research with large language models[J/OL]. Nature, 2023, 624(7992): 570-578. DOI: 10.1038/s41586-023-06792-0.

- [122] HONG S, ZHUGE M, CHEN J, et al. MetaGPT: Meta programming for a multi-agent collaborative framework[C]//The Twelfth International Conference on Learning Representations. 2024.
- [123] JIANG D, REN X, LIN B Y. LLM-blender: Ensembling large language models with pairwise ranking and generative fusion[C/OL]//ROGERS A, BOYD-GRABER J, OKAZAKI N. Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). Toronto, Canada: Association for Computational Linguistics, 2023: 14165-14178. DOI: 10.18653/v1/2023.acl-long.792.
- [124] TEAM K, BAI T, BAI Y, et al. Kimi k2.5: Visual agentic intelligence[Z]. 2026.
- [125] TANG X, ZOU A, ZHANG Z, et al. MedAgents: Large language models as collaborators for zero-shot medical reasoning[C/OL]//KU L W, MARTINS A, SRIKUMAR V. Findings of the Association for Computational Linguistics: ACL 2024. Bangkok, Thailand: Association for Computational Linguistics, 2024: 599-621. DOI: 10.18653/v1/2024.findings-acl.33.
- [126] LIANG T, HE Z, JIAO W, et al. Encouraging divergent thinking in large language models through multi-agent debate[C/OL]//AL-ONAIZAN Y, BANSAL M, CHEN Y N. Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing. Miami, Florida, USA: Association for Computational Linguistics, 2024: 17889-17904. DOI: 10.18653/v1/2024.emnlp-main.992.
- [127] XU Y, WANG S, LI P, et al. Exploring large language models for communication games: An empirical study on werewolf[Z]. 2024.
- [128] CHEN J, SAHA S, BANSAL M. ReConcile: Round-table conference improves reasoning via consensus among diverse LLMs[C/OL]//KU L W, MARTINS A, SRIKUMAR V. Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). Bangkok, Thailand: Association for Computational Linguistics, 2024: 7066-7085. DOI: 10.18653/v1/2024.acl-long.381.
- [129] LIGHT J, CAI M, SHEN S, et al. Avalonbench: Evaluating llms playing the game of avalon[C]//NeurIPS 2023 Foundation Models for Decision Making Workshop. 2023.
- [130] CLAUDE. Orchestrate teams of claude code sessions[EB/OL]. <https://code.claude.com/docs/en/agent-teams>.
- [131] ZHANG J, XIANG J, YU Z, et al. AFlow: Automating agentic workflow generation[C]//The Thirteenth International Conference on Learning Representations. 2025.
- [132] QIAN C, LIU W, LIU H, et al. Chatdev: Communicative agents for software development[C]//Proceedings of the 62nd annual meeting of the association for computational linguistics (volume 1: Long papers). 2024: 15174-15186.
- [133] LIU Z, ZHANG Y, LI P, et al. A dynamic LLM-powered agent network for task-oriented agent collaboration[C]//First Conference on Language Modeling. 2024.
- [134] LI Y, REN S, WU P, et al. Learning distilled collaboration graph for multi-agent perception[C]//NIPS '21: Proceedings of the 35th International Conference on Neural Information Processing Systems. Red Hook, NY, USA: Curran Associates Inc., 2021.
- [135] KIM Y, PARK C, JEONG H, et al. Mdagents: an adaptive collaboration of llms for medical decision-making[C]//NIPS '24: Proceedings of the 38th International Conference on Neural Information Processing Systems. Red Hook, NY, USA: Curran Associates Inc., 2024.
- [136] BROWNE C B, POWLEY E, WHITEHOUSE D, et al. A survey of monte carlo tree search methods[J/OL]. IEEE Transactions on Computational Intelligence and AI in Games, 2012, 4 (1): 1-43. DOI: 10.1109/TCIAIG.2012.2186810.
- [137] ZHENG L, WANG R, WANG X, et al. Synapse: Trajectory-as-exemplar prompting with memory for computer control[C]//The Twelfth International Conference on Learning Representations. 2024.
- [138] HU M, CHEN T, CHEN Q, et al. HiAgent: Hierarchical working memory management for solving long-horizon agent tasks with large language model[C/OL]//CHE W, NABENDE J, SHUTOVA E, et al. Proceedings of the 63rd Annual Meeting of the Association for Computa-

- tional Linguistics (Volume 1: Long Papers). Vienna, Austria: Association for Computational Linguistics, 2025: 32779-32798. DOI: 10.18653/v1/2025.acl-long.1575.
- [139] HONG S, LIN Y, LIU B, et al. Data interpreter: An LLM agent for data science[C/OL]//CHE W, NABENDE J, SHUTOVA E, et al. Findings of the Association for Computational Linguistics: ACL 2025. Vienna, Austria: Association for Computational Linguistics, 2025: 19796-19821. DOI: 10.18653/v1/2025.findings-acl.1016.
 - [140] HU M Y, DURME B V, ANDREAS J, et al. Sample-efficient online learning in LM agents via hindsight trajectory rewriting[C]//NeurIPS 2025 Workshop on Bridging Language, Agent, and World Models for Reasoning and Planning. 2025.
 - [141] WANG Z Z, MAO J, FRIED D, et al. Agent workflow memory[C]//42nd International Conference on Machine Learning. 2025.
 - [142] ZHOU S, ZHOU T, YANG Y, et al. Wall-e 2.0: World alignment by neurosymbolic learning improves world model-based llm agents[Z]. 2025.
 - [143] WANG Z Z, GANDHI A, NEUBIG G, et al. Inducing programmatic skills for agentic tasks[C]//Second Conference on Language Modeling. 2025.
 - [144] QIU J, JUAN X, WANG Y, et al. Agentdistill: Training-free agent distillation with generalizable mcp boxes[Z]. 2025.
 - [145] CAO Z, DENG J, YU L, et al. Remember me, refine me: A dynamic procedural memory framework for experience-driven agent evolution[Z]. 2025.
 - [146] ANDRYCHOWICZ M, WOLSKI F, RAY A, et al. Hindsight experience replay[C]//GUYON I, LUXBURG U V, BENGIO S, et al. Advances in Neural Information Processing Systems: volume 30. Curran Associates, Inc., 2017.
 - [147] KIM N, ONG K T I, HWANG Y, et al. PRINCIPLES: Synthetic strategy memory for proactive dialogue agents[C/OL]//CHRISTODOULOPOULOS C, CHAKRABORTY T, ROSE C, et al. Findings of the Association for Computational Linguistics: EMNLP 2025. Suzhou, China: Association for Computational Linguistics, 2025: 21329-21368. DOI: 10.18653/v1/2025.findings-emnlp.1164.
 - [148] FU Y, KIM D K, KIM J, et al. Autoguide: Automated generation and selection of context-aware guidelines for large language model agents[C]//The Thirty-eighth Annual Conference on Neural Information Processing Systems. 2024.
 - [149] ZHANG G, REN H, ZHAN C, et al. Memevolve: Meta-evolution of agent memory systems[Z]. 2025.
 - [150] XIAO Y, LI Y, WANG H, et al. Toolmem: Enhancing multimodal agents with learnable tool capability memory[J]. arXiv preprint arXiv:2510.06664, 2025.
 - [151] VISHWAKARMA A, LEE H, SURESH M, et al. Cognitive weave: Synthesizing abstracted knowledge with a spatio-temporal resonance graph[Z]. 2025.
 - [152] NAM J, YOON J, CHEN J, et al. Ds-star: Data science agent via iterative planning and verification[Z]. 2025.
 - [153] CAI Z, GUO X, PEI Y, et al. Flex: Continuous agent evolution via forward learning from experience[Z]. 2025.
 - [154] LAURENCEAU J P, BARRETT L F, PIETROMONACO P R. Intimacy as an interpersonal process: the importance of self-disclosure, partner disclosure, and perceived partner responsiveness in interpersonal exchanges.[J]. Journal of personality and social psychology, 1998, 74 (5): 1238.
 - [155] WEI T, SACHDEVA N, COLEMAN B, et al. Evo-memory: Benchmarking llm agent test-time learning with self-evolving memory[Z]. 2025.
 - [156] WU C K, TAM Z R, LIN C Y, et al. Streambench: Towards benchmarking continuous improvement of language agents[C/OL]//GLOBERSON A, MACKAY L, BELGRAVE D, et al. Advances in Neural Information Processing Systems: volume 37. Curran Associates, Inc., 2024: 107039-107063. DOI: 10.52202/079017-3398.
 - [157] WU D, WANG H, YU W, et al. Longmemeval: Benchmarking chat assistants on long-term

- interactive memory[C]//The 13th International Conference on Learning Representations. 2025.
- [158] WAN L, MA W. Storybench: A dynamic benchmark for evaluating long-term memory with multi turns[Z]. 2025.
 - [159] LI X, BANTUPALLI J, DHARMANI R, et al. Toward multi-session personalized conversation: A large-scale dataset and hierarchical tree framework for implicit reasoning[C/OL]//CHRISTODOULOPOULOS C, CHAKRABORTY T, ROSE C, et al. Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing. Suzhou, China: Association for Computational Linguistics, 2025: 11493-11506. DOI: 10.18653/v1/2025.emnlp-main.580.
 - [160] JIANG B, HAO Z, CHO Y M, et al. Know me, respond to me: Benchmarking LLMs for dynamic user profiling and personalized responses at scale[C]//Second Conference on Language Modeling. 2025.
 - [161] XIA M, RÜHLE V, RAJMOHAN S, et al. Minerva: A programmable memory test benchmark for language models[C]//42nd International Conference on Machine Learning. 2025.
 - [162] KIM J, CHAY W, HWANG H, et al. Dialsim: A dialogue simulator for evaluating long-term multi-party dialogue understanding of conversational agents[Z]. 2025.
 - [163] ZHAO S, HONG M, LIU Y, et al. Do LLMs recognize your preferences? evaluating personalized preference following in LLMs[C]//The 13th International Conference on Learning Representations. 2025.
 - [164] ZHOU H, CHEN Y, GUO S, et al. Memento: Fine-tuning llm agents without fine-tuning llms[Z]. 2025.
 - [165] AI Q, TANG Y, WANG C, et al. Memorybench: A benchmark for memory and continual learning in llm systems[Z]. 2025.
 - [166] ZHENG J, CAI X, LI Q, et al. Lifelongagentbench: Evaluating llm agents as lifelong learners[Z]. 2025.
 - [167] TAN H, ZHANG Z, MA C, et al. MemBench: Towards more comprehensive evaluation on the memory of LLM-based agents[C/OL]//CHE W, NABENDE J, SHUTOVA E, et al. Findings of the Association for Computational Linguistics: ACL 2025. Vienna, Austria: Association for Computational Linguistics, 2025: 19336-19352. DOI: 10.18653/v1/2025.findings-acl.989.
 - [168] WU T, CHEN Z, WENG Z, et al. Knowme-bench: Benchmarking person understanding for lifelong digital companions[Z]. 2026.
 - [169] ZHANG Z, DAI Q, CHEN L, et al. Memsim: A bayesian simulator for evaluating memory of llm-based personal assistants[Z]. 2024.
 - [170] DU Y, WANG H, ZHAO Z, et al. PerLTQA: A personal long-term memory dataset for memory classification, retrieval, and fusion in question answering[C]//WONG K F, ZHANG M, XU R, et al. Proceedings of the 10th SIGHAN Workshop on Chinese Language Processing (SIGHAN-10). Bangkok, Thailand: Association for Computational Linguistics, 2024: 152-164.
 - [171] JIA Z, LIU Q, LI H, et al. Evaluating the long-term memory of large language models[C]//Findings of the Association for Computational Linguistics: ACL 2025. 2025: 19759-19777.
 - [172] WANG Q, CHENG Z, ZHANG S, et al. Memgovern: Enhancing code agents through learning from governed human experiences[J]. arXiv preprint arXiv:2601.06789, 2026.
 - [173] SRIVASTAVA S S, HE H. Memorygraft: Persistent compromise of llm agents via poisoned experience retrieval[J]. arXiv preprint arXiv:2512.16962, 2025.
 - [174] SUNIL B D, SINHA I, MAHESHWARI P, et al. Memory poisoning attack and defense on memory based llm-agents[Z]. 2026.
 - [175] ZHAN Q, LIANG Z, YING Z, et al. Injecagent: Benchmarking indirect prompt injections in tool-integrated large language model agents[C]//Findings of the Association for Computational Linguistics ACL 2024. 2024: 10471-10506.
 - [176] DONG S, XU S, HE P, et al. Memory injection attacks on llm agents via query-only interaction[J]. arXiv preprint arXiv:2503.03704, 2025.
 - [177] CHEN Z, XIANG Z, XIAO C, et al. Agentpoison: Red-teaming llm agents via poisoning memory or knowledge bases[J]. Advances in Neural Information Processing Systems, 2024, 37:

- 130185-130213.
- [178] PATLAN A S, SHENG P, HEBBAR S A, et al. Real ai agents with fake memories: Fatal context manipulation attacks on web3 agents[J]. arXiv preprint arXiv:2503.16248, 2025.
 - [179] LIN X, NING Y, ZHANG J, et al. Llm-based agents suffer from hallucinations: A survey of taxonomy, methods, and directions[J]. arXiv preprint arXiv:2509.18970, 2025.
 - [180] WANG B, HE W, ZENG S, et al. Unveiling privacy risks in llm agent memory[C]//Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). 2025: 25241-25260.
 - [181] SHI Z, WAN G, HUANG W, et al. Privacy-enhancing paradigms within federated multi-agent systems[J]. arXiv preprint arXiv:2503.08175, 2025.
 - [182] REZAZADEH A, LI Z, LOU A, et al. Collaborative memory: Multi-user memory sharing in llm agents with dynamic access control[J]. arXiv preprint arXiv:2505.18279, 2025.
 - [183] MCMURRAY C, TIRMAZI H. Whistledown: Combining user-level privacy with conversational coherence in llms[J]. arXiv preprint arXiv:2511.13319, 2025.
 - [184] BARNETT S, KURNIAWAN S, THUDUMU S, et al. Seven failure points when engineering a retrieval augmented generation system[C]//Proceedings of the IEEE/ACM 3rd International Conference on AI Engineering-Software Engineering for AI. 2024: 194-199.
 - [185] ZHANG H, HUANG J, MEI K, et al. Agent security bench (asb): Formalizing and benchmarking attacks and defenses in llm-based agents[J]. arXiv preprint arXiv:2410.02644, 2024.
 - [186] LIN K, SNELL C, WANG Y, et al. Sleep-time compute: Beyond inference scaling at test-time[Z]. 2025.
 - [187] YANG H, XU Z, XIONG F, et al. A first-principles theory of slow thinking and active perception[J/OL]. Preprint, 2026: 1-107. DOI: 10.13140/RG.2.2.21094.28488.